

Rocket: Betting on "What to Build" in the Age of Vibe Coding

1. Executive Summary

Rocket, a startup based in Surat, India, has built its business on a contrarian insight: as AI coding tools like Cursor, Replit, Lovable, and Claude Code make writing software nearly free and instantaneous, the scarce resource shifts from *execution* to *judgment* — knowing what to build in the first place. In April 2026, Rocket launched **Rocket 1.0**, a platform it calls the world's first "Vibe Solutioning" product, which bundles strategic research, product building, and continuous competitive intelligence into a single workflow. The company positions its higher tiers as a low-cost substitute for traditional management consulting, generating "McKinsey-grade" strategy reports from simple text prompts.

This case study examines Rocket's origin as a vibe-coding tool, its strategic pivot upstream into business strategy, its business model and unit economics, and the risks and open questions the model raises.

2. Company Background

Founded by	Vishal Virani (CEO), Rahul Shingala, Deepak Dhanak (COO)
Headquarters	Surat, India (with a US office in Palo Alto)
Team size	~57–58 employees, primarily in Surat
Prior venture	Virani is also CEO of DhiWise, a developer-workflow automation company
Seed funding	\$15 million (September 2025), led by Salesforce Ventures, with Accel and Together Fund participating

Users	Grew from 400,000 (Sept 2025) to over 1.5 million across 180 countries by April 2026

Rocket originally launched in beta as **Rocket.new**, a "vibe-coding" tool competing with Lovable, Cursor, and Bolt. Its early differentiation was a focus on production-readiness rather than throwaway prototypes — what Virani called solving "Day Two" problems (maintenance, iteration, and scaling) rather than just "Day One" (getting a demo working). Roughly 80% of early users were reported to be building substantive applications e-commerce platforms, fintech tools, B2B software, and mental health apps rather than weekend projects.

3. The Strategic Insight: Code Is a Commodity, Judgment Isn't

By early 2026, the competitive landscape for AI-assisted coding had become crowded and largely undifferentiated. As Virani put it, "Everyone can generate the code now... it has become a commodity. But what to build is something which everyone is missing." His framing draws a sharp line between *running a business* and *building a codebase* arguing that the more valuable, less-automated layer of the stack sits **before** code generation begins: figuring out which market to enter, whether an idea holds up, and how to price against competitors.

This is the strategic rationale behind Rocket 1.0: rather than compete purely on coding speed or code quality a race with diminishing differentiation Rocket moved up the value chain into decision support, a space traditionally occupied by strategy consultancies charging thousands of dollars per engagement.

4. The Product: Rocket 1.0

Launched in April 2026, Rocket 1.0 connects three capabilities into one workflow, corresponding to the full lifecycle of a product decision:

1. **Solve (before the build)** — Takes an open business question (e.g., which market to enter, whether an idea is viable, how to price against competitors) and returns a structured, evidence-backed recommendation, delivered as a consulting-style Product Requirement Document (PRD) in PDF format, covering pricing, unit economics, and go-to-market strategy.

2. **Build (during the build)** — Turns the validated strategy directly into a production-ready product (web app, mobile app, dashboard, or customer portal), continuing from the same context rather than requiring the user to switch tools or re-explain the problem.
3. **Intelligence (after launch)** — Continuously monitors competitors across their website, social channels, reviews, job postings, and ad spend, aiming to *interpret* combinations of signals (e.g., a pricing change plus a wave of enterprise hires plus new case studies) as a single strategic narrative rather than isolated data points.

The analytical engine reportedly draws on **more than 1,000 data sources**, including Meta's ad libraries and Similarweb's API, alongside Rocket's own web crawlers.

A caveat surfaced in early testing: journalists who tested the platform before launch found that some outputs appeared to be synthesized from existing public data known pricing models, general user-behavior patterns, and competitive heuristics rather than derived from proprietary or independently verified sources. This means outputs may resemble well-informed synthesis more than bespoke primary research, and Rocket itself acknowledges this by offering human support for users who run into trouble.

5. Business Model and Unit Economics

Rocket monetizes through tiered subscriptions that bundle building, strategy, and intelligence capabilities at increasing price points:

Tier	Price/month	Capability
Entry	\$25	Application building only
Strategy & Research	\$250	~2–3 "McKinsey-grade" strategy reports plus product builds
Full Platform	\$350	Adds competitive intelligence monitoring

Key reported metrics (as of April 2026):

- Annualized average revenue per user (ARPU): roughly \$4,000 for its higher tiers (implying most revenue is concentrated among paying strategy/intelligence customers rather than the \$25 entry tier)
- Gross margins: over 50%
- Customer mix: 20–30% small and medium-sized businesses
- Geographic revenue mix (from the earlier Rocket.new period): US ~26%, Europe ~15–20%, India ~10%

The pricing logic is explicit: a \$250–350/month subscription is framed as a fraction of the cost of traditional strategy consulting, which can run into the thousands of dollars per engagement — making the value proposition especially attractive to solo founders and SMBs who could never previously afford a "Fortune 100-style" strategy team.

6. Competitive Positioning

Rocket sits at the intersection of two categories that have not historically overlapped:

- **Vibe-coding tools** (Cursor, Replit, Lovable, Bolt, Claude Code) — optimized for speed of execution, largely agnostic to whether the thing being built is a good idea.
- **Strategy consulting** (McKinsey-style firms) — optimized for judgment and market insight, but slow, expensive, and disconnected from execution.

By fusing these into one workflow, Rocket is attempting to create a category it calls "Vibe Solutioning," positioning itself not merely as a faster version of either category, but as the connective layer between them take a business question, get a validated strategy, get a shippable product, and get continuous competitive visibility afterward, all without switching tools or re-explaining context to a new system.

This is a deliberate bet that the real value of AI's next wave isn't writing code faster, but shortening the entire loop from idea to informed decision to shipped product to market feedback.

7. Risks and Open Questions

1. **Verifiability of "consulting-grade" claims.** If outputs are partly synthesized from public data patterns rather than proprietary primary research, the "McKinsey-grade" framing may overstate rigor relative to what a human consulting team would deliver — particularly for novel or niche markets where public pricing/behavior data is thin.
 2. **Liability and decision risk.** Strategy documents that inform pricing, market entry, or go-to-market decisions carry higher stakes than a generated app. Errors or hallucinated "facts" in a strategy report could lead a founder to make a costly real-world decision on faulty grounds hence Rocket's own acknowledgment that users should validate outputs.
 3. **Category crowding.** As competitors like Zoho, Emergent, and Composio also expand from coding into surrounding workflows, Rocket's "Solve → Build → Intelligence" loop is a differentiator today but may be replicated quickly.
 4. **Revenue concentration.** With ARPU reportedly averaging around \$4,000/year but an entry tier at just \$25/month, the business likely depends heavily on converting free/low-tier users into the \$250–350 strategy and intelligence tiers — a classic freemium upsell challenge.
 5. **Trust in AI-generated strategic judgment.** The product's core promise (offloading "what to build" decisions to AI) is a higher-trust ask than "help me write this function." Whether users are willing to make consequential business bets based on AI-synthesized research, rather than treating it as a first-pass input to human judgment, remains an open market question.
-

8. Key Takeaways

- **Strategic sequencing matters more than technical novelty.** Rocket's edge isn't a superior LLM it's recognizing where in the product-development value chain AI has *not yet* commoditized judgment, and moving there before competitors do.
- **Bundling adjacent workflows can justify premium pricing.** By connecting strategy, building, and competitive intelligence in one context-aware system, Rocket converts what would be three separate purchases (a consultant, a coding tool, a competitive-intelligence subscription) into one higher-priced bundle.
- **AI-native founders come from unconventional geographies.** A team built in Surat outside India's traditional tech hubs reached 1.5 million users across 180 countries within roughly a year of its initial beta launch, illustrating how AI-native tooling can compress the geographic advantages that used to matter in enterprise software and consulting.

- **The line between "helpful synthesis" and "verified research" is the central tension of this business model**, and how Rocket manages that gap through disclaimers, human support, or improved sourcing will likely determine whether it is trusted for genuinely high-stakes decisions or treated as a fast first draft.

Sources: TechCrunch (April 2026); IT Voice (April 2026); IndexBox (April 2026); AIBoomi (October 2025); Analytics India Magazine (August 2025).